

Chapter 1

Algorithmic Fundamentals

1.1 Environment Detection

1.1.1 Image Processing^[ME]

The first step towards a desired steering angle is processing the image. The ideal result of the image processing would be an image with perfectly applied edge detection, where all the edges are one colour and the rest is another. We decided to turn all edges white and everything else black.

We started out by using OpenCV to convert our colorspace image into grayscale. By using the *threshold* function from OpenCV it was easy to filter out a certain spectrum of brightness in the image and drawing contours around the remaining edges was done with a mixture of *findContours* and *drawContours*.

This resulted in an image that was only black with white contours around objects, but due to how the *threshold* function was built it was rather hard to achieve a consistent good quality image because of the different lightning conditions in the room. To counter this problem we tried to create an adaptive threshold using the light sensor on the front facing ultrasonic sensor, this did increase the quality of the image. But while testing this we stumbled upon OpenCV's *Canny* function.

Canny is a function specifically made for edge detection, it uses the *Canny86* algorithm (See Algorithmic References for more information about *Canny*). Using this we get an acceptable image to work with at almost all times. We noticed that the algorithm itself might be a bit resource heavy but not so heavy that it would pose as a problem, especially not when we only analyze half the image as explained in the Implementation section [JONAS MAKE LINK].

1.1.2 Lane Detection^[DK]

When the image is processed with clear edges it is time to apply lane detection logic. Every iteration we loop through pixels horizontally and vertically, looking for white or light-grey pixels. For the left and right lane markings, the pixel we start looking from is vertically a pixel very close to the bottom of the picture, we refer to it as our control scanline. Horizontally, we start from the middle of the picture. We then loop through the pixels to the left and right of the starting point, looking for white pixels. The number of pixels

looped through before finding the desired pixel is then stored and used as a distance measurement when applying the lane following.

The stop-line detection works the same way as the detection of the left and right lane markings. When looking for stop-lines however, we have two separate starting points, they are offset ± 50 pixels to the left and right of the starting point we used in the lane marking detection. We then loop through pixels on the vertically instead of horizontally. The reason we check the stopline at two separate locations is to simply increase robustness, more specifically in this case these two distances are compared to check whether the line detected is more or less horizontal. Additional robustness in this detection is achieved by checking this for a few iterations before actually treating it as a stop-line. If these robustness checks are passed, we forward the distance to the stop-line to the DecisionMaker, in which we adjust the speed appropriately. Before this is done however, we ran a few tests and concluded that there was no point in sending any stop-line distances if the line was too far away. Infact, we ended up only sending the distance if it was low enough that the DecisionMaker would consider stopping. This was not the initial idea, at first we wanted to send farther distances and have a certain distance at which the car slows down until it reaches the threshold of stopping. However, seeing as we had some issues getting the car to run at higher speeds with our current implementation, we felt that there was no purpose in slowing down due to the car's speed already being low enough to simply stop. The car also stopped very smoothly when setting the speed to 0.

Furthermore, when no lane markings are found on either side during the lane-detection, a flag we call quality is set to false. What this means is that the data from the LaneFollower is not to be trusted when it comes to decision making. The intention was that this value were to be handled in the DecisionMaker, potentially lowering the speed and perhaps even setting the steering wheel angle to 0. We unfortunately did not have time to handle this in the DecisionMaker although it was handled in the Lane Follower.




Figure 1.1: LaneFollower screenshot

1.1.3 Ultrasonic Sensors^[JE]

The ultrasonic sensors use the IC protocol for communication. As discussed further in the Hardware-Software Integration section **[JONAS MAKE LINK]**, this link seems sensitive to fluctuations in voltage, and although we reduced this sensitivity by mounting external pull-up resistors connected to the I²C bus, the sensors can still respond with unexpected 0's occasionally instead of the current distance measured. Therefore, we discard any readings of 0 until it occurs at least three consecutive times.

All values are averaged with a circular buffer, currently having the size of 4 elements. These mechanisms may reduce reaction time slightly for sudden obstacles, but we considered the reduction in noise to be worth this, taking into account which challenges the miniature car was expected to handle.

Values are capped to 90 cm (anything above this will still be sent as 90), because values above this will likely be unusable, noisy fluctuations not di-

rectly interesting to the current state of the car.

The light sensor of the front ultrasonic sensor is also read, and used to determine the strength of the headlights. Earlier we discussed the possibility of adjusting the camera image processing threshold values based on the light sensor reading, which is why the light sensor reading is also transmitted to the high-level board.

1.1.4 Infrared Sensors^[JE]

The infrared sensors are read using the analog-to-digital converters of the STM32. We based our formula for translating sensor voltage levels to distance on a formula designed for Arduinos, which have a lower ADC resolution. To compensate for the difference in resolution, we simply bitshift the values.

To minimize noise, each time we read the sensor values we let the ADC read 4 samples for each sensor. Then we translate the average sampled voltage level for each sensor into a distance value, which is in turn averaged with the previous distance value.

Values are capped to 28 cm, because any values above this will be unusable, noisy fluctuations.

1.1.5 Wheel Encoder^[TDJ]

The wheel encoder works by measuring wheel revolutions using infrared reflections. In a proper setup, it should be possible to not only measure distance traveled, but speed as well. However, fitting the wheel encoder properly was hard and the results were at times unpredictable.

The circumference of the wheel measured was 20.42 cm. This wheel had nine stripes on its interiors, which reflected the infrared light back to the sensor. After a meter of travel, the sensor should have measured 44 stripes. These measurements could be used to calculate both speed and distance traveled using the amount of stripes per meter.

Initially, the wheel encoder ran within its own thread. The rate of reliable results was not satisfactory with this setup, so it was setup to rely on interrupt routines which were triggered when a new value was read. This increased the reliability of results quite significantly, but was still insufficient at operating

speeds with an inaccuracy of 10 - 15%

The particular wheel encoder used is clearly specified by the manufacturer as only intended for a specific wheel (different from the wheels used in this project). If we would have had more time, we would consider using for example a computer mouse to measure distance travelled.

1.2 Vehicle Control^[JE]

1.2.1 Safety Mechanism

If no valid control data is received through the serial connection for a given time, the wheels are centered and the engine is stopped. This constitutes a basic safety mechanism for a broken connection or execution problems on the high-level board.

1.3 Steering Control

The steering wheel angle is forwarded agnostically from the high-level board to the servo. In the OpenDaVinci components, radians are used. In the proxy, these are converted to degrees, centered around 90 degrees to make sure we can use unsigned bytes for transmitting. Hence 90 degrees means forward, 60 means full left turn and so on. We roughly see 30 degrees or about 0.5 radians as the maximum turning angle in each direction.

1.3.1 Engine Control

For controlling the engine power, we use a set of prefixed pulsewidth values. The high-level board sends an integer number, where 0 means reverse, 1 means neutral, 2 means forward slowly, and 3 means cruise speed forward. Speeds above 1.0 in the OpenDaVINCI-based components are seen as cruise speed. The conversion occurs in the proxy component. The solution of indexed speed modes has safety benefits, and indeed we did not have the problems of unexpected engine power surges that several other groups did.

We do not use any speed regulator mechanism to adjust the engine power based on how the current speed relates to the desired speed, and therefore had to manually adjust the PWM values to use the car with low battery power. We deemed the speed measurement from the wheel encoder to be too

unreliable to use as a basis for adjusting the engine power. If we would have had the time to explore more reliable ways to measure speed or distance, a speed regulator could have been helpful.

1.3.2 Light Control

Using the neopixel LED strips, we gave the car headlights, tail lights, brake lights, indicator lights, reverse light and an RC mode light (blue). The reverse light, indicator lights (flashing orange) and brake lights are controlled by the components running on the high-level board. The RC mode light is automatic. The tail lights are always lit, and the headlights have a strength that adapts to the surrounding light levels (as measured by the light sensor in the front ultrasonic sensor), which can help reduce the cameras exposure time in dark environments.